

Interpretable Card-Play Strategies from Evolved Logic Networks

A Neurosymbolic Agent for Sueca

André Cardoso

June 2026

DETI, Universidade de Aveiro

Most strong card-playing AIs are black boxes: can competitive play be human-auditable?

The status quo

Deep PIMC and neural policy nets.

Millions of floating-point weights →
opaque.

“Why did you play that card?”

“... the matrix says so.” “... I saw 14 million futures, and this is the only one where we win.”

Our goal

Evolved logic gates.


A handful of **human-readable IF/THEN rules.**

“I duck because the trick is cheap and I want to build a void.”

Thesis: discrete logic gates + neuroevolution → competitive play *and* interpretability.

Sueca's distinctive rules create a rich strategic texture

Table layout (4 players, 2 teams)




● Team 0&2 ● Team 1&3

Key rules

- **40 cards** (no 8,9,10)
- Rank: **A > 7 > K > J > Q > 6-2**
- 7 = *manilha* (2nd-highest!)
- **120 points/deal**; game-point tiers
- Must **follow suit**; trump cuts
- **Public void tracking**
- Imperfect information: you see only your hand + history

These rules define the strategic vocabulary the WANN must learn.

The agent reads a public-information belief state through evolved logic gates and outputs a strategic intent



35 belief features
Hand, trick state, voids,
boss detection, "can I win?"

3 intents
MAX_FORCE,
EFFICIENT_WIN,
EQUITY_BUILDER

Styled resolver
Every intent $\approx$ Elite;
100% legal outputs

Weight-agnostic networks replace learned weights with evolved logic-gate topology

Node types

Aggregation	Activation
SUM Σx_i	IDENTITY x
MIN = AND	NOT $1-x$
MAX = OR	THRESHOLD $x > 0.5$

- No MEAN (precision), no SIGMOID (rules)
- All outputs clamped to $[0, 1]$

Why it compiles to rules

- Connections: **sign only** (± 1)
- Shared weight sweep:
 $W \in \{-2, -1, -0.5, 0.5, 1, 2\}$
- Topology must work at *all* scales
- Every gate is a logic primitive
→ **Boolean-ish rules**
- Tie-breaking: **random** among max intents

The architecture constraint is the interpretability guarantee.

Three play intents: styled deviations around a strong Elite policy

MAX_FORCE

Elite + **control**.

When trump-long: lead a **low trump** to draw opponents' high trumps.

Otherwise: play **aggressively**.

"I have trump length — let me thin their trumps."

EFFICIENT_WIN

Exactly Elite.

Play the **cheapest** card that wins the trick.

If you can't win: cheapest legal sacrifice.

"Win cheap, lose cheap."

EQUITY_BUILDER

Elite + **tempo**.

Lead **shortest** side-suit (build a void for later cuts).


Duck cheap tricks (≤ 2 pts) when not last to play.

"I'll save my points and build void threats."

Each intent shares Elite's core; each benchmarks $\approx$ Elite individually.

The WANN's job is to choose **when** to deviate, not **how** to play.

Phase 0: supervised bootstrap on PIMC expert data, starting from zero connections



PFS-NEAT from **zero** connections. Dataset split by `Am_I_Leading` into lead/follow.
Adaptive 2-stage mutation filter.

Phase 1: co-evolutionary self-play with dynamic per-decision brain routing



Lead and follow brains co-evolve. Dynamic routing via `Am_I_Leading` flag.

Delta-fitness with CRN. HOF + MAP-Elites opponent pools.

A cheap supra-Elite teacher made the champion $4.5\times$ simpler at equal strength

Problem diagnosed

- Deep PIMC **only tied** Elite
- Root cause: **myopic leaf eval** (raw score, no positional estimate)
- Depth doesn't fix a blind leaf

Solution

- Flat Monte-Carlo PIMC + **Elite** **playouts**
- Rollout policy improvement:
1-ply + Elite $\geq$ Elite
- **62% vs. Elite**, $\sim 1000\times$ faster

What the better teacher bought


	Earlier	Final
vs. Elite	52.7%	52.1%
Hidden gates	132	29
Connections	188	49
Teacher	$\alpha\text{-}\beta$	rollout

Statistically identical strength $\rightarrow 4.5\times$

simpler.

The ceiling is the resolver, not the teacher.

The champion beats Elite 52.1% $\pm$ 1.8% ($n = 3000$)



Tournament matrix (win%, $n = 3000$)

	Rnd	Old	Elite	WANN
Random	50.0	10.8	4.7	5.1
Old	89.2	50.0	32.5	32.3
Elite	95.3	67.5	50.0	47.9
WANN	95.0	67.7	52.1	50.0

- 95% CI [50.3%, 53.9%] **excludes 50%**
- Card points: 60.2 vs. 59.8
- WANN vs. Old: 67.7% (Elite: 67.5%)
- Prior best **30.2%**; first time the project beat Elite

Three root-cause fixes turned a collapsed network into an Elite-beating champion



Problem

Network collapsed to
“always EFFICIENT_WIN”
Intents unequally weak
20% / 37% / 25%

Fix #1

Styled resolver (Fix 1)
Every intent $\approx$ Elite
48% / 47% / 46%
Collapse $\rightarrow$ merely ties
Result: 52.7% vs. Elite
(132 gates, 188 conns)

Fix #2 & #3

Rollout teacher (Fix 2+3)
62% vs. Elite (itself)
 $\rightarrow$ 4.5 $\times$ simpler champion
Result: 52.1% vs. Elite
(29 gates, 49 conns)
equal strength

The follow brain compiles to 5 logic gates at depth 5 (verified behavior-preserving)

Compiled FOLLOW-brain rules (after constant-folding + alias inlining)

```
hidden_42 = NOT(Holds_Boss_Led)
hidden_48 = (1 + Any_Opp_Void_Led)
hidden_41 = NOT((1 + h_42 + 2*hidden_48))
hidden_40 = THRESHOLD((h_41 + Trump_Count) > 0.5)
hidden_55 = THRESHOLD(hidden_40 > 0.5)
MAX_FORCE      = 2*Game_Pts_Remaining
EFFICIENT_WIN   = 0
EQUITY_BUILDER = 1 + 2*hidden_41 + hidden_55
```


LEAD brain (10 gates, depth 6):

```
EQUITY = THRESHOLD(NOT(Trump_Count))
MAX_FORCE = (ahead & point-dense
             & partner-winning)
           OR (no short side suit)
```

EFFICIENT_WIN = 0 in both brains: the net never delegates to pure Elite; it always applies one of the two overlays.

Verified on 2000 random states: the removed symbols carry zero information.

Before folding: the raw evolved follow-brain topology



17 hidden gates, 44 connections. Constant-folding and alias inlining collapse the *active* logic to the

5 gates on the previous slide, a $4.5\times$ reduction versus the earlier champion.

The styled resolver enabled the win but also sets the ceiling: we measured four strength levels

	Pure intent	Champion	Oracle	RF predictor
vs. Elite	47–48%	52.1%	58%	63–71%

Pure intent: each of the 3 intents played alone (no WANN)

→ $\approx$ Elite the “raised floor” that prevents collapse.

Oracle: perfect intent selection at every decision

→ 58% is the **absolute cap** of the current vocabulary.

RF = Random Forest trained on belief features to predict the oracle’s choice.

63–71% accuracy shows the **features carry usable signal** the WANN has not yet learned.

The +2% gap (47% floor → 52% champion) is the **learned deviation overlay**.

What limits us


- 3 intents nearly equivalent in fitness
→ weak learning signal
- Lead decision harder: RF gap +2.5

pp₁ (vs. follow +7.9 pp)

What could raise the cap

- More-differentiated intents
- Opponent-card inference features
- Knowledge distillation

You can explore the champion's decision-making interactively in the browser



React + WASM: Rust game engine compiled to WebAssembly
Play Sueca against the WANN, watch the logic flow in real time.

sueca.mycsina.xyz

Richer intents and opponent-card inference are the path to breaking the 58% ceiling

Future directions

1. **Richer intents**
more-differentiated play styles
→ raise the oracle cap above 58%
2. **Opponent-card inference**
can a logic circuit learn to infer hidden cards from presumed optimal play?
3. **Knowledge distillation**
champion → even smaller student
→ Pareto curve of strength vs. size
4. **Collaboration**
can a logic circuit display emergent collaboration capabilities in Sueca?

What we showed

- Evolved an **interpretable**, Elite-beating Sueca agent
- Verified pipeline:
logic gates → **5 readable rules**
- Built an interactive website that allows to play and explore the agent in a real game environment
- Created a Monte Carlo tree search algorithm to estimate the true potential of any given game position.
- All code, data, and rules are reproducible

Competitive, interpretable card-play AI is possible: the rules fit on one slide.

Thank you. Questions?

Key references

Gaier & Ha 2019	Weight-Agnostic Neural Networks
Stanley & Miikkulainen 2002	NEAT
Mouret & Clune 2015	MAP-Elites
Reisinger et al. 2004	Coevolving modular (L-NEAT)

Reproducibility

All commands in README.md

Canonical champion: `checkpoints/production/2026-06-14-2/
cargo build -p sueca_wann -release && cargo test -all`

André Cardoso
DETI, Universidade de Aveiro