

Interpretable Card-Play Strategies from Evolved Logic Networks

A Neurosymbolic Agent for Sueca

André Cardoso

June 2026

DETI, Universidade de Aveiro

Most strong card-playing AIs are black boxes: can competitive play be human-auditable?

The status quo

Deep PIMC and neural policy nets.

Millions of floating-point weights →
opaque.

"Why did you play that card?"

"... the matrix says so."

Our goal

Evolved logic gates.


A handful of **human-readable IF/THEN rules.**

"I duck because the trick is cheap and I want to build a void."

Thesis: discrete logic gates + neuroevolution → competitive play *and* interpretability.

Sueca's distinctive rules create a rich strategic texture

Table layout (4 players, 2 teams)




Key rules

- **40 cards** (no 8,9,10)
- Rank: **A > 7 > K > J > Q > 6-2**
- 7 = *manilha* (2nd-highest!)
- **120 points/deal**; game-point tiers
- Must **follow suit**; trump cuts
- **Public void tracking**
- Imperfect information: you see only your hand + history

These rules define the strategic vocabulary the WANN must learn.

The agent reads a public-information belief state through evolved logic gates and outputs a strategic intent



35 belief features
Hand, trick state, voids,
boss detection, "can I win?"

3 intents
MAX_FORCE,
EFFICIENT_WIN,
EQUITY_BUILDER

Styled resolver
Every intent $\approx$ Elite;
100% legal outputs

Weight-agnostic networks replace learned weights with evolved logic-gate topology

Node types

Aggregation	Activation
SUM Σx_i	IDENTITY x
MIN = AND	NOT $1-x$
MAX = OR	THRESHOLD $x > 0.5$


- No MEAN (precision), no SIGMOID (rules)
- All outputs clamped to $[0, 1]$

Why it compiles to rules

- Connections: **sign only** (± 1)
- Shared weight sweep:
 $W \in \{-2, -1, -0.5, 0.5, 1, 2\}$
- Topology must work at *all* scales
- Every gate is a logic primitive
→ **Boolean-ish rules**
- Tie-breaking: **random** among max intents

The architecture constraint is the interpretability guarantee.

Training bootstraps on expert demonstrations then switches to co-evolutionary self-play



Phase 0 (gens 0–149)

Supervised pretraining on PIMC data
PFS-NEAT from **zero** connections
Split: *lead* vs. *follow* datasets
Adaptive 2-stage mutation filter

Phase 1 (gens 150–599)

Co-evolutionary self-play
Dynamic per-decision brain routing
Delta-fitness with CRN
HOF + MAP-Elites opponent pools

A cheap supra-Elite teacher made the champion $4.5\times$ simpler at equal strength

Problem diagnosed

- Deep PIMC **only tied** Elite
- Root cause: **myopic leaf eval** (raw score, no positional estimate)
- Depth doesn't fix a blind leaf

Solution

- Flat Monte-Carlo PIMC + **Elite playouts**
- Rollout policy improvement: 1-ply + Elite $\geq$ Elite
- **62% vs. Elite**, $\sim 1000\times$ faster


What the better teacher bought

	Earlier	Final
vs. Elite	52.7%	52.1%
Hidden gates	132	29
Connections	188	49
Teacher	α - β	rollout

Stronger teacher $\rightarrow$ simpler champion.

The ceiling is the resolver, not the teacher.

The champion beats Elite 52.1% $\pm$ 1.8% ($n = 3000$)



Tournament matrix (win%, $n = 3000$)

	Rnd	Old	Elite	WANN
Random	50.0	10.8	4.7	5.1
Old	89.2	50.0	32.5	32.3
Elite	95.3	67.5	50.0	47.9
WANN	95.0	67.7	52.1	50.0

- 95% CI [50.3%, 53.9%] **excludes 50%**
- Card points: 60.2 vs. 59.8
- WANN vs. Old: 67.7% (Elite: 67.5%)
- Prior best **30.2%**; first time the project beat Elite

Three root-cause fixes turned a collapsed network into an Elite-beating champion



Problem

Network collapsed to
“always EFFICIENT_WIN”
Intents unequally weak
20% / 37% / 25%

Fix #1

Styled resolver
Every intent $\approx$ Elite
48% / 47% / 46%
Collapse $\rightarrow$ merely ties

Fix #2 & #3

Best-of-3-intent labeling
Rollout teacher
62% vs. Elite
 $\rightarrow 4.5\times$ simpler

Diagnosed with a training-free harness in seconds, not blind retraining in hours.

The follow brain compiles to 5 logic gates at depth 5 (verified behavior-preserving)

Compiled FOLLOW-brain rules (after constant-folding + alias inlining)

```
hidden_42 = NOT(Holds_Boss_Led)
hidden_48 = (1 + Any_Opp_Void_Led)
hidden_41 = NOT((1 + h_42 + 2*hidden_48))
hidden_40 = THRESHOLD((h_41 + Trump_Count) > 0.5)
hidden_55 = THRESHOLD(hidden_40 > 0.5)
MAX_FORCE      = 2*Game_Pts_Remaining
EFFICIENT_WIN   = 0
EQUITY_BUILDER = 1 + 2*hidden_41 + hidden_55
```


LEAD brain (10 gates, depth 6):

```
EQUITY = THRESHOLD(NOT(Trump_Count))
MAX_FORCE = (ahead & point-dense
             & partner-winning)
           OR (no short side suit)
```

EFFICIENT_WIN = 0 in both brains: the net never delegates to pure Elite; it always applies one of the two overlays.

Verified on 2000 random states: the removed symbols carry zero information.

Before folding: the raw evolved follow-brain topology



17 hidden gates, 44 connections. Constant-folding and alias inlining collapse the *active* logic to the

5 gates on the previous slide, a $4.5\times$ reduction versus the earlier champion.

The styled resolver that enabled the win also sets the strength ceiling, and we measured it

	Any pure intent	Final champion	3-intent oracle	RF ceiling
vs. Elite	47–48%	52.1%	58%	63–71%

Elite floor → every intent $\approx$ Elite individually → collapse merely ties
Oracle cap → perfect intent selection $\approx$ 58% → $\sim$ 8 pp headroom left
+2% = the learned deviation overlay on top of the fixed resolver base.

What limits us


- 3 intents nearly equivalent in fitness
- Weak gradient for learned selection
- Lead decision harder: RF +2.5 pp
- Follow learnable: RF +7.9 pp

What could raise the cap

- More-differentiated intents
- Opponent-card inference features
- Knowledge distillation
- Richer belief state

We measure and report our ceiling; +2% is the learned overlay.

You can explore the champion's decision-making interactively in the browser



React + WASM: Rust game engine compiled to WebAssembly
Play Sueca against the WANN, watch the logic flow in real time.

Richer intents and opponent-card inference are the path to breaking the 58% ceiling

Future directions

1. **Richer intents**
more-differentiated play styles
→ raise the oracle cap above 58%
2. **Opponent-card inference**
can a logic circuit learn to infer hidden cards from optimal play?
3. **Knowledge distillation**
champion → even smaller student
→ Pareto curve of strength vs. size

What we showed

- Evolved an **interpretable**, Elite-beating Sueca agent
- Verified pipeline:
logic gates → **5 readable rules**
- Honest analysis of the resolver-floor ceiling
- Supra-Elite rollout teacher
($\sim 1000\times$ cheaper)
- All code, data, and rules are reproducible

Competitive, interpretable card-play AI is possible: the rules fit on one slide.

Thank you. Questions?

Key references

Gaier & Ha 2019	Weight-Agnostic Neural Networks
Stanley & Miikkulainen 2002	NEAT
Mouret & Clune 2015	MAP-Elites
Reisinger et al. 2004	Coevolving modular (L-NEAT)

Reproducibility

All commands in README.md

```
Canonical champion: checkpoints/production/2026-06-14-2/  
cargo build -p sueca_wann -release && cargo test -all
```

André Cardoso
DETI, Universidade de Aveiro